

---

## Main Goals of the Project








### 1. Real-World Health Problem Understanding

- Predict whether a person is likely to have **diabetes** based on health-related data like:
  - Blood sugar levels
  - BMI
  - Insulin levels
  - Age
  - Pregnancy count
- Helps simulate how **AI models are used in medical diagnostics** in hospitals or health tech apps.

---

### 2. Learn and Practice ML Workflow with TensorFlow

You will experience the **complete ML pipeline**:

Phase	Goal
 <b>Data Collection</b>	Load real-world health dataset (Pima Indians)
 <b>Data Analysis</b>	Use Pandas, Seaborn to understand the data
 <b>Preprocessing</b>	Normalize features using StandardScaler
 <b>Model Building</b>	Build a neural network with TensorFlow
 <b>Model Training</b>	Train using real data to find patterns
 <b>Evaluation</b>	Test model accuracy and compare predictions
 <b>Visualization</b>	Plot loss curves, correlation heatmap

---

### 3. Apply Deep Learning to a Classification Task

- This is a **binary classification** problem:
    - Outcome = 1: Person has diabetes
    - Outcome = 0: Person does not have diabetes
  - Goal: **Minimize error** and correctly classify new patients
-





## 4. Model Interpretation and Insight

By the end of the project, you'll be able to:

- Read model accuracy and loss
- Understand how changing features affects predictions
- Think about feature importance in healthcare diagnostics




---

### Bonus Learning Objectives

Skill	Description
 <b>Model design</b>	Learn structure: Input → Hidden Layers → Output
 <b>Interpret loss/accuracy</b>	Understand performance with validation loss
 <b>Visualization skills</b>	Build useful plots to explain findings
 <b>Explainability</b>	Learn how to present your model in a healthcare context

---

### Real-World Use Cases

-  Used by **doctors** to assist in screening patients
  -  Integrated in **mobile health apps** like Apple Health, Practo, etc.
  -  Helps **insurance companies** assess risk of applicants
- 

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns #multiple var. in ur dataset ,statistical high level ,quick
import tensorflow as tf #dev. by google , deep learning,intelligent system
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

#url

url = "https://raw.githubusercontent.com/jbrownlee/Datasets/refs/heads/master/pima-indians-diabetes.data.csv"

columns =
['Pregnancies','Glucose','BloodPressure','SkinThickness','Insulin','BMI','DiabetesPedigreeFunction',
Age,'Outcome']
```

```
df=pd.read_csv(url,header =None,names=columns)
print(df.head())
print(df.shape)
print(df.describe())
print(df.info())
```

```
#heatmap
plt.figure(figsize=(10,6))
sns.heatmap(df.corr(),annot=True,cmap='coolwarm')
plt.title('Diabetes Outcome Distribution')
plt.show()

#diabetes count
sns.countplot(x='Outcome',data=df)
plt.title('Count of Outcome')
plt.xlabel('Outcome')
plt.ylabel('Count')
plt.show()
```

```
#prepare the data
X = df.drop('Outcome',axis=1) #ip
y = df['Outcome'] #op
print(X.head())
print(y.head())
```

```
#scale the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X) #mean = 0 , sd =1
print(X_scaled)
```

```
#split the data
```

```
X_train,X_test,y_train,y_test = train_test_split(X_scaled,y,test_size=0.2,random_state=42) # 20%
testing data 80% training data
```

```
#build model (fully connected (layer ip , hidden layer,output = perceptron)
model = tf.keras.Sequential([
    tf.keras.layers.Dense(64,activation='relu',input_shape=(X_train.shape[1],)),
    tf.keras.layers.Dense(32,activation='relu'),
    tf.keras.layers.Dense(1,activation='sigmoid')
])
model.compile(optimizer='adam',loss='binary_crossentropy',metrics=['accuracy'])
```

```
# Assign the return value of model.fit to the history variable
history = model.fit(X_train,y_train,epochs=100,batch_size=32,validation_data=(X_test,y_test))
```

```
#evaluate the model
loss,accuracy = model.evaluate(X_test,y_test)
print(f'Test Loss:{loss:.4f}')
print(f'Test Accuracy:{accuracy:.4f}')
```

```
# Print the keys in the history object to identify the correct key for accuracy
print(history.history.keys())
```

```
# Use the correct key for accuracy, which is likely 'accuracy' based on the metrics=['accuracy']
# If the print statement above shows a different key (e.g., 'acc'), replace 'accuracy' with that key
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train','Validation'],loc='upper left')
plt.show()
```

```
y_pred = model.predict(X_test)
print("\n Predictions vs Actual:")
for i in range(10):
    print(f'Predicted:{y_pred[i][0]:.2f},Actual:{y_test.iloc[i]}')
```

-----